

Experiment-4

Name: -G. Akshya

Regno: -192425132

Course Code: -MLA0305

Slot: B

Title: -Exploration vs. Exploitation Analysis Using Multi-Armed Bandit Algorithms

Aim- To analyze the exploration-exploitation trade-off using a Multi-Armed Bandit problem and understand how an agent selects between exploring new actions and exploiting the best-known action.

Procedure

1. Create a Multi-Armed Bandit environment with multiple arms.
2. Assign different reward probabilities to the arms.
3. Initialize the estimated reward for every arm.
4. Select an exploration-exploitation strategy such as ϵ -greedy.
5. Generate a random value to determine whether the agent explores or exploits.
6. During exploration, select an action randomly.
7. During exploitation, select the action having the highest estimated reward.
8. Execute the selected action and obtain its reward.
9. Update the estimated reward of the selected arm.
10. Repeat the process for multiple trials.
11. Calculate cumulative rewards and average rewards.
12. Plot the performance of the different actions.

Source Code

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.mixture import GaussianMixture
from IPython.display import display, FileLink
from google.colab import files

np.random.seed(42)

print("=" * 80)
print("EXPERIMENT 4")
print("EXPLORATION VS. EXPLOITATION ANALYSIS USING MULTI-ARMED BANDITS")
print("=" * 80)

# =====
# 1. GMM DATASET
# =====

N = 500
```

```

base_data = np.random.uniform(
    0.1,
    0.9,
    size=(N, 4)
)

gmm = GaussianMixture(
    n_components=3,
    random_state=42
)

gmm.fit(base_data)

generated = gmm.sample(N)[0]

dataset = pd.DataFrame(
    generated,
    columns=[
        "Arm_1_Probability",
        "Arm_2_Probability",
        "Arm_3_Probability",
        "Arm_4_Probability"
    ]
)

dataset = dataset.clip(0.05, 0.95).round(3)

# =====
# 2. DISPLAY 10 SAMPLES
# =====

print("\n10 SAMPLE DATASET RECORDS")

display(dataset.head(10))

print("Total Records:", len(dataset))

# =====
# 3. DATASET SUMMARY
# =====

summary_dataset = dataset.describe().T[
    ["min", "max", "mean", "std"]
].round(3)

```

```

print("\nDATASET SUMMARY")

display(summary_dataset)

# =====
# 4. DOWNLOAD
# =====

filename = "Experiment_4_Bandit_GMM_Dataset.csv"

dataset.to_csv(
    filename,
    index=False
)

print("\nDOWNLOAD COMPLETE DATASET:")

display(FileLink(filename))

files.download(filename)

# =====
# 5. BANDIT PARAMETERS
# =====

TRUE_PROBABILITIES = np.array([
    0.20,
    0.40,
    0.65,
    0.85
])

ARMS = 4
ROUNDS = 500

# =====
# 6. EPSILON GREEDY
# =====

def epsilon_greedy(epsilon=0.10):

    counts = np.zeros(ARMS)
    rewards = np.zeros(ARMS)

    reward_history = []

```

```

selected_arms = []

for t in range(ROUNDS):

    if t < ARMS:
        arm = t

    elif np.random.rand() < epsilon:
        arm = np.random.randint(ARMS)

    else:

        means = np.divide(
            rewards,
            counts,
            out=np.zeros(ARMS),
            where=counts != 0
        )

        arm = np.argmax(means)

    reward = int(
        np.random.rand()
        < TRUE_PROBABILITIES[arm]
    )

    counts[arm] += 1
    rewards[arm] += reward

    reward_history.append(reward)
    selected_arms.append(arm)

    return (
        np.array(reward_history),
        np.array(selected_arms),
        counts
    )

# =====
# 7. UCB
# =====

def ucb():

    counts = np.zeros(ARMS)

```

```

rewards = np.zeros(ARMS)

reward_history = []
selected_arms = []

for t in range(ROUNDS):

    if t < ARMS:

        arm = t

    else:

        means = rewards / counts

        confidence = np.sqrt(
            2 * np.log(t + 1)
            / counts
        )

        arm = np.argmax(
            means + confidence
        )

        reward = int(
            np.random.rand()
            < TRUE_PROBABILITIES[arm]
        )

        counts[arm] += 1
        rewards[arm] += reward

        reward_history.append(reward)
        selected_arms.append(arm)

    return (
        np.array(reward_history),
        np.array(selected_arms),
        counts
    )

# =====
# 8. THOMPSON SAMPLING
# =====

```

```

def thompson_sampling():

    alpha = np.ones(ARMS)
    beta = np.ones(ARMS)

    reward_history = []
    selected_arms = []

    for _ in range(ROUNDS):

        samples = np.random.beta(
            alpha,
            beta
        )

        arm = np.argmax(samples)

        reward = int(
            np.random.rand()
            < TRUE_PROBABILITIES[arm]
        )

        if reward == 1:
            alpha[arm] += 1
        else:
            beta[arm] += 1

        reward_history.append(reward)
        selected_arms.append(arm)

    counts = np.bincount(
        selected_arms,
        minlength=ARMS
    )

    return (
        np.array(reward_history),
        np.array(selected_arms),
        counts
    )

# =====
# 9. RUN ALGORITHMS
# =====

```

```

eps_rewards, eps_arms, eps_counts = epsilon_greedy()
ucb_rewards, ucb_arms, ucb_counts = ucb()
ts_rewards, ts_arms, ts_counts = thompson_sampling()

# =====
# 10. CUMULATIVE REWARD
# =====

eps_cumulative = np.cumsum(eps_rewards)
ucb_cumulative = np.cumsum(ucb_rewards)
ts_cumulative = np.cumsum(ts_rewards)

# =====
# 11. REGRET
# =====

optimal_probability = max(TRUE_PROBABILITIES)

eps_regret = np.cumsum(
    optimal_probability - TRUE_PROBABILITIES[eps_arms]
)

ucb_regret = np.cumsum(
    optimal_probability - TRUE_PROBABILITIES[ucb_arms]
)

ts_regret = np.cumsum(
    optimal_probability - TRUE_PROBABILITIES[ts_arms]
)

# =====
# 12. COMPARISON TABLE
# =====

comparison = pd.DataFrame({
    "Algorithm": [
        "Epsilon-Greedy",
        "UCB",
        "Thompson Sampling"
    ],
    "Total Reward": [
        eps_rewards.sum(),
        ucb_rewards.sum(),
        ts_rewards.sum()
    ]
})

```

```

],

    "Average Reward": [
        eps_rewards.mean(),
        ucb_rewards.mean(),
        ts_rewards.mean()
    ],

    "Final Regret": [
        eps_regret[-1],
        ucb_regret[-1],
        ts_regret[-1]
    ],

    "Most Selected Arm": [
        np.argmax(eps_counts) + 1,
        np.argmax(ucb_counts) + 1,
        np.argmax(ts_counts) + 1
    ]
}).round(3)

print("\n" + "=" * 80)
print("ALGORITHM COMPARISON")
print("=" * 80)

display(comparison)

# =====
# 13. SUMMARY TABLE
# =====

best_algorithm = comparison.loc[
    comparison["Average Reward"].idxmax(),
    "Algorithm"
]

summary = pd.DataFrame({
    "Metric": [
        "Dataset Records",
        "Number of Arms",
        "Number of Rounds",
        "Best Arm",
        "Best Arm Probability",
        "Best Performing Algorithm"
    ],

```



```

        "Result": [
            len(dataset),
            ARMS,
            ROUNDS,
            np.argmax(TRUE_PROBABILITIES) + 1,
            max(TRUE_PROBABILITIES),
            best_algorithm
        ]
    })

print("\nSUMMARY RESULTS")

display(summary)

# =====
# 14. VISUALIZATION 1
# ARM PROBABILITIES
# =====

plt.figure(figsize=(10, 6))

plt.bar(
    ["Arm 1", "Arm 2", "Arm 3", "Arm 4"],
    TRUE_PROBABILITIES,
    color=[
        "#ff6b6b",
        "#4dabf7",
        "#51cf66",
        "#ffd43b"
    ],
    edgecolor="black"
)

plt.title(
    "True Reward Probability of Each Bandit Arm",
    fontsize=16,
    fontweight="bold"
)

plt.ylabel("Probability")
plt.ylim(0, 1)

plt.show()

```

```

# =====
# 15. VISUALIZATION 2
# CUMULATIVE REWARD
# =====

plt.figure(figsize=(11, 6))

plt.plot(
    eps_cumulative,
    label="Epsilon-Greedy",
    color="#ff6b6b",
    linewidth=3
)

plt.plot(
    ucb_cumulative,
    label="UCB",
    color="#4dabf7",
    linewidth=3
)

plt.plot(
    ts_cumulative,
    label="Thompson Sampling",
    color="#51cf66",
    linewidth=3
)

plt.title(
    "Cumulative Reward Comparison",
    fontsize=16,
    fontweight="bold"
)

plt.xlabel("Round")
plt.ylabel("Cumulative Reward")

plt.legend()
plt.grid(alpha=0.25)

plt.show()

# =====
# 16. VISUALIZATION 3
# CUMULATIVE REGRET

```

```

# =====

plt.figure(figsize=(11, 6))

plt.plot(
    eps_regret,
    label="Epsilon-Greedy",
    color="#ff6b6b",
    linewidth=3
)

plt.plot(
    ucb_regret,
    label="UCB",
    color="#4dabf7",
    linewidth=3
)

plt.plot(
    ts_regret,
    label="Thompson Sampling",
    color="#51cf66",
    linewidth=3
)

plt.title(
    "Cumulative Regret Comparison",
    fontsize=16,
    fontweight="bold"
)

plt.xlabel("Round")
plt.ylabel("Cumulative Regret")

plt.legend()
plt.grid(alpha=0.25)

plt.show()

# =====
# 17. VISUALIZATION 4
# ARM SELECTION FREQUENCY
# =====

x = np.arange(ARMS)

```

```
width = 0.25

plt.figure(figsize=(11, 6))

plt.bar(
    x - width,
    eps_counts,
    width,
    label="Epsilon-Greedy",
    color="#ff6b6b"
)

plt.bar(
    x,
    ucb_counts,
    width,
    label="UCB",
    color="#4dabf7"
)

plt.bar(
    x + width,
    ts_counts,
    width,
    label="Thompson Sampling",
    color="#51cf66"
)

plt.xticks(
    x,
    ["Arm 1", "Arm 2", "Arm 3", "Arm 4"]
)

plt.title(
    "Arm Selection Frequency",
    fontsize=16,
    fontweight="bold"
)

plt.ylabel("Number of Selections")
plt.legend()

plt.show()

# =====
```

```

# 18. VISUALIZATION 5
# AVERAGE REWARD
# =====

plt.figure(figsize=(10, 6))

plt.bar(
    comparison["Algorithm"],
    comparison["Average Reward"],
    color=[
        "#ff6b6b",
        "#4dabf7",
        "#51cf66"
    ],
    edgecolor="black"
)

plt.title(
    "Average Reward Comparison",
    fontsize=16,
    fontweight="bold"
)

plt.ylabel("Average Reward")
plt.xticks(rotation=15)

plt.show()

print("\n" + "=" * 80)
print("FINAL RESULT")
print("=" * 80)

print(
    f"Best performing algorithm: {best_algorithm}"
)

```

10 SAMPLE DATASET RECORDS

	Arm_1_Probability	Arm_2_Probability	Arm_3_Probability	Arm_4_Probability
0	0.527	0.450	0.114	0.697
1	0.934	0.546	0.611	0.756
2	0.611	0.480	0.941	0.734
3	0.875	0.692	0.766	0.451
4	0.639	0.787	0.592	0.691
5	0.686	0.328	0.738	0.747
6	0.788	0.063	0.499	0.759
7	0.913	0.797	0.377	0.763
8	0.757	0.338	0.769	0.536
9	0.641	0.553	0.567	0.891

FIG: - 1 SAMPLE DATA

DATASET SUMMARY

	min	max	mean	std
Arm_1_Probability	0.05	0.95	0.480	0.232
Arm_2_Probability	0.05	0.95	0.508	0.222
Arm_3_Probability	0.05	0.95	0.483	0.220
Arm_4_Probability	0.05	0.95	0.515	0.224

FIG: - 2 DATA SET SUMMARY

=====

ALGORITHM COMPARISON

=====

	Algorithm	Total Reward	Average Reward	Final Regret	Most Selected Arm
0	Epsilon-Greedy	404	0.808	20.80	4
1	UCB	387	0.774	45.05	4
2	Thompson Sampling	406	0.812	6.50	4

FIG: - 3 ALGORITHM COMPARISON

Results: -

SUMMARY RESULTS

	Metric	Result
0	Dataset Records	500
1	Number of Arms	4
2	Number of Rounds	500
3	Best Arm	4
4	Best Arm Probability	0.85
5	Best Performing Algorithm	Thompson Sampling

FIG: - 4 SUMMARY RESULTS




FIG: - 5 REWARD PROBABILITY




FIG: - 6 CUMULATIVE REWARD




FIG: - 7 CUMULATIVE REGRET




FIG: - 8 ARM SELECTION




FIG: - 9 AVERAGE REWARD COMPARISON

Result

The exploration vs. exploitation trade-off was successfully analyzed using the Multi-Armed Bandit algorithm. The agent learned to exploit actions with higher expected rewards while still exploring other actions to discover potentially better choices.